


In coroutines, a *flow* is a type that can emit multiple values sequentially, as opposed to *suspend functions* that return only a single value. For example, you can use a flow to receive live updates from a database.

Flows are built on top of coroutines and can provide multiple values. A flow is conceptually a *stream of data* that can be computed asynchronously. The emitted values must be of the same type. For example, a `Flow<Int>` is a flow that emits integer values.

A flow is very similar to an `Iterator` that produces a sequence of values, but it uses suspend functions to produce and consume values asynchronously. This means, for example, that the flow can safely make a network request to produce the next value without blocking the main thread.

There are three entities involved in streams of data:

- A **producer** produces data that is added to the stream. Thanks to coroutines, flows can also produce data asynchronously.
- **(Optional) Intermediaries** can modify each value emitted into the stream or the stream itself.
- A **consumer** consumes the values from the stream.



**Figure 1.** Entities involved in streams of data: consumer, optional intermediaries, and producer.

In Android, a [repository](#) is typically a producer of UI data that has the user interface (UI) as the consumer that ultimately displays the data. Other times, the UI layer is a producer of user input events and other layers of the hierarchy consume them. Layers in between the producer and consumer usually act as intermediaries that modify the stream of data to adjust it to the requirements of the following layer.

## Creating a flow

To create flows, use the [flow builder](#) APIs. The `flow` builder function creates a new flow where you can manually emit new values into the stream of data using the `emit` function.

In the following example, a data source fetches the latest news automatically at a fixed interval. As a suspend function cannot return multiple consecutive values, the data source creates and returns a flow to fulfill this requirement. In this case, the data source acts as the producer.

```
class NewsRemoteDataSource(  
    private val newsApi: NewsApi,  
    private val refreshIntervalMs: Long = 5000  
) {  
    val latestNews: Flow<List<ArticleHeadline>> = flow {  
        while(true) {  
            val latestNews = newsApi.fetchLatestNews()  
            emit(latestNews) // Emits the result of the request to the flow  
            delay(refreshIntervalMs) // Suspends the coroutine for some time  
        }  
    }  
}  
  
// Interface that provides a way to make network requests with suspend functions  
interface NewsApi {  
    suspend fun fetchLatestNews(): List<ArticleHeadline>  
}
```

## Modifying the stream

Intermediaries can use *intermediate operators* to modify the stream of data without consuming the values. These operators are functions that, when applied to a stream of data, set up a chain of operations that aren't executed until the values are consumed in the future. Learn more about intermediate operators in the [Flow reference documentation](#).

In the example below, the repository layer uses the intermediate operator `map` to transform the data to be displayed on the `View`:

```
class NewsRepository(  
    private val newsRemoteDataSource: NewsRemoteDataSource,  
    private val userData: UserData  
) {  
    /**  
     * Returns the favorite latest news applying transformations on the flow.  
     * These operations are lazy and don't trigger the flow. They just transform  
     * the current value emitted by the flow at that point in time.  
     */  
    val favoriteLatestNews: Flow<List<ArticleHeadline>> =  
        newsRemoteDataSource.latestNews
```

```

        // Intermediate operation to filter the list of favorite topics
        .map { news -> news.filter { userData.isFavoriteTopic(it) } }
        // Intermediate operation to save the latest news in the cache
        .onEach { news -> saveInCache(news) }
    }
}

```

Intermediate operators can be applied one after the other, forming a chain of operations that are executed lazily when an item is emitted into the flow. Note that simply applying an intermediate operator to a stream does not start the flow collection.

Use a *terminal operator* to trigger the flow to start listening for values. To get all the values in the stream as they're emitted, use `collect`. You can learn more about terminal operators in the [official flow documentation](#).

As `collect` is a suspend function, it needs to be executed within a coroutine. It takes a lambda as a parameter that is called on every new value. Since it's a suspend function, the coroutine that calls `collect` may suspend until the flow is closed.

Continuing the previous example, here's a simple implementation of a `ViewModel` consuming the data from the repository layer:

```

class LatestNewsViewModel(
    private val newsRepository: NewsRepository
) : ViewModel() {

    init {
        viewModelScope.launch {
            // Trigger the flow and consume its elements using collect
            newsRepository.favoriteLatestNews.collect { favoriteNews ->
                // Update View with the latest favorite news
            }
        }
    }
}

```

Collecting the flow triggers the producer that refreshes the latest news and emits the result of the network request on a fixed interval. As the producer remains always active with the `while(true)` loop, the stream of data will be closed when the `ViewModel` is cleared and `viewModelScope` is cancelled.

Flow collection can stop for the following reasons:

- The coroutine that collects is cancelled, as shown in the previous example. This also stops the underlying producer.
- The producer finishes emitting items. In this case, the stream of data is closed and the coroutine that called `collect` resumes execution.

Flows are *cold* and *lazy* unless specified with other intermediate operators. This means that the producer code is executed each time a terminal operator is called on the flow. In the previous example, having multiple flow collectors causes the data source to fetch the latest news multiple times on different fixed intervals. To optimize and share a flow when multiple consumers collect at the same time, use the [shareIn](#) operator.

## Catching unexpected exceptions

The implementation of the producer can come from a third party library. This means that it can throw unexpected exceptions. To handle these exceptions, use the [catch](#) intermediate operator.

```
class LatestNewsViewModel(  
    private val newsRepository: NewsRepository  
) : ViewModel() {  
  
    init {  
        viewModelScope.launch {  
            newsRepository.favoriteLatestNews  
                // Intermediate catch operator. If an exception is thrown,  
                // catch and update the UI  
                .catch { exception -> notifyError(exception) }  
                .collect { favoriteNews ->  
                    // Update View with the latest favorite news  
                }  
        }  
    }  
}
```

[catch](#) can also [emit](#) items to the flow. The example repository layer could [emit](#) the cached values instead:

```
class NewsRepository(...) {  
    val favoriteLatestNews: Flow<List<ArticleHeadline>> =  
        newsRemoteDataSource.latestNews  
            .map { news -> news.filter { userData.isFavoriteTopic(it) } }  
            .onEach { news -> saveInCache(news) }  
            // If an error happens, emit the last cached values  
            .catch { exception -> emit(lastCachedNews()) }  
}
```

In this example, when an exception occurs, the [collect](#) lambda is called, as a new item has been emitted to the stream because of the exception.

## Executing in a different CoroutineContext

By default, the producer of a `flow` builder executes in the `CoroutineContext` of the coroutine that collects from it, and as previously mentioned, it cannot `emit` values from a different `CoroutineContext`. This behavior might be undesirable in some cases. For instance, in the examples used throughout this topic, the repository layer shouldn't be performing operations on `Dispatchers.Main` that is used by `viewModelScope`.

To change the `CoroutineContext` of a flow, use the intermediate operator `flowOn`. `flowOn` changes the `CoroutineContext` of the *upstream flow*, meaning the producer and any intermediate operators applied *before (or above)* `flowOn`. The *downstream flow* (the intermediate operators *after* `flowOn` along with the consumer) is not affected and executes on the `CoroutineContext` used to `collect` from the flow. If there are multiple `flowOn` operators, each one changes the upstream from its current location.

```
class NewsRepository(
    private val newsRemoteDataSource: NewsRemoteDataSource,
    private val userData: UserData,
    private val defaultDispatcher: CoroutineDispatcher
){
    val favoriteLatestNews: Flow<List<ArticleHeadline>> =
        newsRemoteDataSource.latestNews
            .map { news -> // Executes on the default dispatcher
                news.filter { userData.isFavoriteTopic(it) }
            }
            .onEach { news -> // Executes on the default dispatcher
                saveInCache(news)
            }
            // flowOn affects the upstream flow ↑
            .flowOn(defaultDispatcher)
            // the downstream flow ↓ is not affected
            .catch { exception -> // Executes in the consumer's context
                emit(lastCachedNews())
            }
}
```

With this code, the `onEach` and `map` operators use the `defaultDispatcher`, whereas the `catch` operator and the consumer are executed on `Dispatchers.Main` used by `viewModelScope`.

As the data source layer is doing I/O work, you should use a dispatcher that is optimized for I/O operations:

```
class NewsRemoteDataSource(
    ...,
    private val ioDispatcher: CoroutineDispatcher
){
    val latestNews: Flow<List<ArticleHeadline>> = flow {
        // Executes on the IO dispatcher
    }
```

```
    ...  
  }  
  .flowOn(ioDispatcher)  
}
```

## Flows in Jetpack libraries

Flow is integrated into many Jetpack libraries, and it's popular among Android third party libraries. Flow is a great fit for live data updates and endless streams of data.

You can use [Flow with Room](#) to be notified of changes in a database. When using [data access objects \(DAO\)](#), return a **Flow** type to get live updates.

```
@Dao  
abstract class ExampleDao {  
    @Query("SELECT * FROM Example")  
    abstract fun getExamples(): Flow<List<Example>>  
}
```

Every time there's a change in the **Example** table, a new list is emitted with the new items in the database.